NAME

llamafile-imatrix — importance matrix builder

SYNOPSIS

`llamafile-imatrix` [flags...] `-m` *`model.gguf`* `-f` *`training.data`* [`-o` *`imatrix.dat`*]

DESCRIPTION

`llamafile-imatrix` Compute an importance matrix for a model and given text dataset. Can be used during quantization to enchance the quality of the quantum models. More information is available here: https://github.com/ggerganov/llama.cpp/pull/4861

OPTIONS

The following options are available:

`--version`
: Print version and exit.

`-h`, `--help`
: Show help message and exit.

`-m` *`FNAME`*, `--model` *`FNAME`*
: Model path in the GGUF file format.

  Default: *models/7B/ggml-model-f16.gguf*

`-f` *`FNAME`*, `--file` *`FNAME`*
: Mandatory path of file containing training data, e.g. *wiki.train.raw*

`-o` *`FNAME`*, `--output-file` *`FNAME`*
: The name of the file where the computed data will be stored. If this flag is missing then *imatrix.dat* is used.

`-ofreq`, `--output-frequency`
: Specifies how often the so far computed result is saved to disk. The default is is 10 (i.e., every 10 chunks).

`-ow`, `--output-weight`
: Specifies if data will be collected for the *output.weight* tensor. Experience indicates that it is better to not utilize the importance matrix when quantizing *output.weight*, so this is set to false by default.

PROTIPS

For faster computation, pass the `-ngl` *`9999`* flag for GPU of floading.

SEE ALSO

llamafile(1), *llamafile-quantize*(1)